

STAT GR5242 Advanced Machine Learning

Improved Residual Network for Image Classification

Author: Jie Li, Xiaofan Zhang, Zhaoyang Wang

UNI: j15246, xz2735, zw2551

```
In [1]: import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
import skimage.transform
import time, os, datetime
from Model import Residual_Unit
from Model import Attention_Block
```

```
In [2]: from tensorflow.keras.datasets import cifar10
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.optimizers import Adam, SGD
from tensorflow.keras.callbacks import ReduceLROnPlateau, EarlyStopping, TensorBoard, LearningRateScheduler
from tensorflow.keras.models import load_model
from tensorflow.keras.layers import Input
from tensorflow.keras.regularizers import l2
from tensorflow.keras.layers import BatchNormalization
from tensorflow.keras.layers import Conv2D
from tensorflow.keras.layers import MaxPooling2D
from tensorflow.keras.layers import Activation
from tensorflow.keras.layers import AveragePooling2D
from tensorflow.keras.layers import Flatten
from tensorflow.keras.layers import Dropout
from tensorflow.keras.layers import Dense
from tensorflow.keras.models import Model
```

```
In [3]: print("TF version: ", tf.__version__)
print("Keras version:", tf.keras.__version__)
```

TF version: 2.0.0
Keras version: 2.2.4-tf

Load CIFAR-10 Data

```
In [4]: # Load the CIFAR10 data.
(x_train, y_train), (x_test, y_test) = cifar10.load_data()
x_train = x_train[:30000, :, :, :]
y_train = y_train[:30000]
x_val = x_train[-5000:, :, :, :]
y_val = y_train[-5000:]

print('x_train shape:', x_train.shape)
print('y_train shape:', y_train.shape)
print('x_validation shape:', x_val.shape)
print('y_validation shape:', y_val.shape)
print('x_test shape:', x_test.shape)
print('y_test shape:', y_test.shape)

# Convert class vectors to binary class matrices.
y_train = to_categorical(y_train, 10)
y_val = to_categorical(y_val, 10)
y_test = to_categorical(y_test, 10)
```

```
x_train shape: (30000, 32, 32, 3)
y_train shape: (30000, 1)
x_validation shape: (5000, 32, 32, 3)
y_validation shape: (5000, 1)
x_test shape: (10000, 32, 32, 3)
y_test shape: (10000, 1)
```

Data Augmentation

```
In [5]: # define generators for training and validation data
train_datagen = ImageDataGenerator(
    featurewise_center=True,
    featurewise_std_normalization=True,
    rotation_range=20,
    width_shift_range=0.2,
    height_shift_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True,
    validation_split=0.2)

val_datagen = ImageDataGenerator(
    featurewise_center=True,
    featurewise_std_normalization=True)

test_datagen = ImageDataGenerator(
    featurewise_center=True,
    featurewise_std_normalization=True)

# compute quantities required for feature normalization
train_datagen.fit(x_train)
val_datagen.fit(x_val)
test_datagen.fit(x_test)
```

Construct Residual Attention Network

Stacking two Attention modules follow by Residual Network

```

In [6]: # define learning rate scheduler
def lr_schedule(epoch):
    lr = 1e-4
    if epoch > 50:
        lr *= 1e-2
    elif epoch > 20:
        lr *= 1e-1
    print('Learning rate:', lr)
    return lr
lr_scheduler = LearningRateScheduler(lr_schedule)

# Residual Attention Network
def AttentionResNet56_mini(shape, in_channel, kernel_size, n_classes, dropout=None, regul

    """
    :param shape: The tuple of input data.
    :param in_channel: The 4-th dimension (channel number) of input weight matrix. For exa
    :param kernel_size: Integer. the shape of the kernel. For example, default kernel_size
    :param n_classes: Integer. The number of target classes. For example, n_classes = 10 m
    :param dropout: Float between 0 and 1. Fraction of the input units to drop.
    :param regularization: Float. Fraction of the input units to drop.
    """

    input_data = Input(shape=shape) # 32x32x32
    x = Conv2D(in_channel, kernel_size=kernel_size, padding='same')(input_data) # 32x32x32
    x = BatchNormalization()(x)
    x = Activation('relu')(x)
    x = MaxPooling2D(pool_size=2, padding='same')(x) # 16x16x32

    out_channel = in_channel * 4
    x = Residual_Unit(x, in_channel, out_channel) # 16x16x128
    x = Attention_Block(x, skip=1)

    in_channel = out_channel // 2
    out_channel = in_channel * 4
    x = Residual_Unit(x, in_channel, out_channel, stride=2) # 8x8x256
    x = Attention_Block(x, skip=1)

    in_channel = out_channel // 2
    out_channel = in_channel * 4
    x = Residual_Unit(x, in_channel, out_channel, stride=1) # 4x4x1024
    x = Residual_Unit(x, in_channel, out_channel)
    x = Residual_Unit(x, in_channel, out_channel)

    x = AveragePooling2D(pool_size=4, strides=1)(x) # 1x1x1024
    x = Flatten()(x)

    output = Dense(n_classes, activation='softmax')(x)
    model = Model(input_data, output)

    return model

def training(model, log_name, batch_size=128, epc=60):
    batch_size = batch_size
    epc = epc

```

```
start = time.time()

# define training generator
train_generator = train_datagen.flow(x_train, y_train, batch_size=batch_size)
step_size_train = train_generator.n // train_generator.batch_size

# define validation generator
val_generator = val_datagen.flow(x_val, y_val, batch_size=batch_size)
step_size_val = val_generator.n // val_generator.batch_size

# define test validation generator
test_generator = test_datagen.flow(x_test, y_test, batch_size=batch_size)
step_size_test = test_generator.n // test_generator.batch_size

# usefull callbacks
log_dir='Logs/' + log_name
tensorboard_callback = TensorBoard(log_dir=log_dir, histogram_freq=1)
lr_reducer = ReduceLROnPlateau(monitor='val_accuracy', factor=0.1, patience=5, verbose=1)
early_stopper = EarlyStopping(monitor='val_accuracy', patience=15, verbose=1)

model.fit_generator(train_generator,
                    steps_per_epoch = step_size_train,
                    epochs = epc,
                    validation_data = val_generator,
                    validation_steps = step_size_val,
                    callbacks=[tensorboard_callback, lr_reducer, lr_scheduler, early_stopper])

end = time.time()
print("Time taken by above cell is {}".format((end-start)/60))

# evaluation
val_scores = model.evaluate_generator(val_generator, verbose=0)
test_scores = model.evaluate_generator(test_generator, verbose=1)
print('validation loss:', val_scores[0])
print('validation accuracy:', val_scores[1])
print('Test loss:', test_scores[0])
print('Test accuracy:', test_scores[1])

return model
```

```
In [24]: # define model
model = AttentionResNet56_mini(shape=(32, 32, 3), in_channel=32, kernel_size=5, n_classes=10)

# define loss, metrics, optimizer
optimizer = SGD(lr = lr_schedule(0), momentum=0.9, nesterov=True)
# optimizer = Adam(lr = lr_schedule(0) )
model.compile(optimizer, loss='categorical_crossentropy', metrics=['accuracy'])
model.summary()
```

batch_normalization_144 (BatchN	(None, 4, 4, 256)	1024	conv2d_199[0][0]
<hr/>			
activation_149 (Activation)	(None, 4, 4, 256)	0	batch_normalization_144[0][0]
<hr/>			
conv2d_200 (Conv2D)	(None, 4, 4, 256)	65792	activation_149[0][0]
<hr/>			
conv2d_201 (Conv2D)	(None, 4, 4, 256)	65792	max_pooling2d_9[0][0]
<hr/>			
add_46 (Add)	(None, 4, 4, 256)	0	conv2d_200[0][0] conv2d_201[0][0]
<hr/>			
batch_normalization_145 (BatchN	(None, 4, 4, 256)	1024	add_46[0][0]

Training and Evaluation

Train our base model with **nesterov SGD** optimizer

```
In [28]: # training
model = training(model, '56mini-SGD-Base')

Epoch 35/60
234/234 [=====] - 75s 322ms/step - loss: 1.2155 - accuracy:
0.5676 - val_loss: 1.0880 - val_accuracy: 0.6100
Learning rate: 1e-05
Epoch 36/60
234/234 [=====] - 76s 325ms/step - loss: 1.2221 - accuracy:
0.5666 - val_loss: 1.0958 - val_accuracy: 0.6118
Learning rate: 1e-05
Epoch 37/60
234/234 [=====] - 76s 326ms/step - loss: 1.2130 - accuracy:
0.5657 - val_loss: 1.0982 - val_accuracy: 0.6158
Learning rate: 1e-05
Epoch 38/60
234/234 [=====] - 75s 321ms/step - loss: 1.2162 - accuracy:
0.5680 - val_loss: 1.1053 - val_accuracy: 0.6142
Learning rate: 1e-05
Epoch 39/60
233/234 [=====>.] - ETA: 0s - loss: 1.2119 - accuracy: 0.5691
Epoch 00039: ReduceLROnPlateau reducing learning rate to 9.99999747378752e-07.
234/234 [=====] - 74s 316ms/step - loss: 1.2123 - accuracy:
```

Instead of using SGD, change to **Adam** optimizer

```
In [29]: # define model
model = AttentionResNet56_mini(shape=(32,32,3), in_channel=32, kernel_size=5, n_classes=10)

# define loss, metrics, optimizer
optimizer = Adam(lr = lr_schedule(0) )
model.compile(optimizer, loss='categorical_crossentropy', metrics=['accuracy'])
```

Learning rate: 0.0001

```
In [30]: # training
model = training(model, '56mini-Adam-Base')
0.0000 - val_loss: 0.0000 - val_accuracy: 0.0000
Learning rate: 1.0000000000000002e-06
Epoch 58/60
234/234 [=====] - 75s 321ms/step - loss: 0.4818 - accuracy:
0.8298 - val_loss: 0.3698 - val_accuracy: 0.8686
Learning rate: 1.0000000000000002e-06
Epoch 59/60
234/234 [=====] - 74s 318ms/step - loss: 0.4814 - accuracy:
0.8320 - val_loss: 0.3681 - val_accuracy: 0.8694
Learning rate: 1.0000000000000002e-06
Epoch 60/60
234/234 [=====] - 75s 319ms/step - loss: 0.4861 - accuracy:
0.8299 - val_loss: 0.3607 - val_accuracy: 0.8694
Time taken by above cell is 76.77415573596954.
79/79 [=====] - 9s 112ms/step - loss: 0.6643 - accuracy: 0.7
854
validation loss: 0.35428522787988187
validation accuracy: 0.8696
Test loss: 0.6642644971231871
Test accuracy: 0.7854
```

Improvements

Since the overfitting issue occurs, we would like to add regularization methods to generalize our model as following:

- Batch normalization
- L2 norm regularization
- Dropout
- Early stop


```

In [7]: def AttentionResNet56_mini(shape, in_channel, kernel_size, n_classes, dropout=None, regul

        """
        :param shape: The tuple of input data.
        :param in_channel: The 4-th dimension (channel number) of input weight matrix. For exa
        :param kernel_size: Integer. the shape of the kernel. For example, default kernel_size
        :param n_classes: Integer. The number of target classes. For example, n_classes = 10 m
        :param dropout: Float between 0 and 1. Fraction of the input units to drop.
        :param regularization: Float. Fraction of the input units to drop.
        """

        input_data = Input(shape=shape) # 32x32x32
        x = Conv2D(in_channel, kernel_size=kernel_size, padding='same')(input_data) # 32x32x32
        x = BatchNormalization()(x)
        x = Activation('relu')(x)
        x = MaxPooling2D(pool_size=2, padding='same')(x) # 16x16x32

        out_channel = in_channel * 4
        x = Residual_Unit(x, in_channel, out_channel) # 16x16x128
        x = Attention_Block(x, skip=1)

        in_channel = out_channel // 2
        out_channel = in_channel * 4
        x = Residual_Unit(x, in_channel, out_channel, stride=2) # 8x8x256
        x = Attention_Block(x, skip=1)

        in_channel = out_channel // 2
        out_channel = in_channel * 4
        x = Residual_Unit(x, in_channel, out_channel, stride=1) # 4x4x1024
        x = Residual_Unit(x, in_channel, out_channel)
        x = Residual_Unit(x, in_channel, out_channel)

        # add BN and Activation
        x = BatchNormalization()(x) # new
        x = Activation('relu')(x) # new
        x = AveragePooling2D(pool_size=4, strides=1)(x) # 1x1x1024
        x = Flatten()(x)

        if dropout:
            x = Dropout(dropout)(x) # new
        output = Dense(n_classes, kernel_regularizer=l2(regularization), activation='softmax')
        model = Model(input_data, output)

        return model

```

```
In [8]: # define model
model = AttentionResNet56_mini(shape=(32, 32, 3), in_channel=32, kernel_size=5, n_classes=100)

# define loss, metrics, optimizer
optimizer = Adam(lr = lr_schedule(0) )
model.compile(optimizer, loss='categorical_crossentropy', metrics=['accuracy'])
model.summary()
```

conv2d_50 (Conv2D)	(None, 4, 4, 256)	65792	max_pooling2d_2[0][0]
add_11 (Add)	(None, 4, 4, 256)	0	conv2d_49[0][0] conv2d_50[0][0]
batch_normalization_37 (Batch Normalization)	(None, 4, 4, 256)	1024	add_11[0][0]
activation_38 (Activation)	(None, 4, 4, 256)	0	batch_normalization_37[0][0]
conv2d_51 (Conv2D)	(None, 4, 4, 256)	65792	activation_38[0][0]

```
In [34]: # training
model = training(model, '56mini-regularization')
0.0000 - val_loss: 0.3535 - val_accuracy: 0.8912
Learning rate: 1.0000000000000002e-06
Epoch 58/60
234/234 [=====] - 77s 330ms/step - loss: 0.4481 - accuracy:
0.8562 - val_loss: 0.3535 - val_accuracy: 0.8912
Learning rate: 1.0000000000000002e-06
Epoch 59/60
234/234 [=====] - 76s 325ms/step - loss: 0.4526 - accuracy:
0.8545 - val_loss: 0.3472 - val_accuracy: 0.8942
Learning rate: 1.0000000000000002e-06
Epoch 60/60
234/234 [=====] - 77s 329ms/step - loss: 0.4478 - accuracy:
0.8544 - val_loss: 0.3475 - val_accuracy: 0.8932
Time taken by above cell is 77.04772863785426.
79/79 [=====] - 9s 114ms/step - loss: 0.6473 - accuracy: 0.8
029
validation loss: 0.34524615183472634
validation accuracy: 0.893
Test loss: 0.6473370205752457
Test accuracy: 0.8029
```

Reduce the batch size and increase the epoch size

```
In [9]: def AttentionResNet56_mini(shape, in_channel, kernel_size, n_classes, dropout=None, regul

    """
    :param shape: The tuple of input data.
    :param in_channel: The 4-th dimension (channel number) of input weight matrix. For exa
    :param kernel_size: Integer. the shape of the kernel. For example, default kernel_size
    :param n_classes: Integer. The number of target classes. For example, n_classes = 10 m
    :param dropout: Float between 0 and 1. Fraction of the input units to drop.
    :param regularization: Float. Fraction of the input units to drop.
    """

    input_data = Input(shape=shape) # 32x32x32
    x = Conv2D(in_channel, kernel_size=kernel_size, padding='same')(input_data) # 32x32x32
    x = MaxPooling2D(pool_size=2, padding='same')(x) # 16x16x32

    out_channel = in_channel * 4
    x = Residual_Unit(x, in_channel, out_channel) # 16x16x128
    x = Attention_Block(x, skip=2)

    in_channel = out_channel // 2
    out_channel = in_channel * 4
    x = Residual_Unit(x, in_channel, out_channel, stride=2) # 8x8x256
    x = Attention_Block(x, skip=1)
    x = Attention_Block(x, skip=1)

    in_channel = out_channel // 2
    out_channel = in_channel * 4
    x = Residual_Unit(x, in_channel, out_channel, stride=1) # 4x4x1024
    x = Residual_Unit(x, in_channel, out_channel)
    x = Residual_Unit(x, in_channel, out_channel)

    # add BN and Activation
    x = BatchNormalization()(x) # new
    x = Activation('relu')(x) # new
    x = AveragePooling2D(pool_size=4, strides=1)(x) # 1x1x1024
    x = Flatten()(x)

    if dropout:
        x = Dropout(dropout)(x) # new
    output = Dense(n_classes, kernel_regularizer=l2(regularization), activation='softmax')
    model = Model(input_data, output)

    return model
```

```
In [10]: # define model
model = AttentionResNet56_mini(shape=(32, 32, 3), in_channel=32, kernel_size=5, n_classes=10)

# define loss, metrics, optimizer
optimizer = Adam(lr = lr_schedule(0) )
model.compile(optimizer, loss='categorical_crossentropy', metrics=['accuracy'])
model.summary()
```

conv2d_141 (Conv2D)	(None, 8, 8, 256)	65792	conv2d_140[0][0]
activation_97 (Activation)	(None, 8, 8, 256)	0	batch_normalization_94[0][0]
activation_104 (Activation)	(None, 8, 8, 256)	0	conv2d_141[0][0]
conv2d_130 (Conv2D)	(None, 8, 8, 256)	65792	activation_97[0][0]
conv2d_131 (Conv2D)	(None, 8, 8, 256)	65792	add_29[0][0]
lambda_3 (Lambda)	(None, 8, 8, 256)	0	activation_104[0][0]
add_31 (Add)	(None, 8, 8, 256)	0	conv2d_130[0][0]

```
In [36]: # training
model = training(model, '56paper', batch_size=64, epc=80)
```

468/468 [=====] - 220s 470ms/step - loss: 0.4368 - accuracy: 0.8500 - val_loss: 0.3108 - val_accuracy: 0.8992
Learning rate: 1.0000000000000002e-06
Epoch 58/80
468/468 [=====] - 220s 471ms/step - loss: 0.4320 - accuracy: 0.8539 - val_loss: 0.3067 - val_accuracy: 0.9008
Learning rate: 1.0000000000000002e-06
Epoch 59/80
468/468 [=====] - 218s 467ms/step - loss: 0.4306 - accuracy: 0.8537 - val_loss: 0.3079 - val_accuracy: 0.8986
Learning rate: 1.0000000000000002e-06
Epoch 60/80
468/468 [=====] - 220s 470ms/step - loss: 0.4278 - accuracy: 0.8553 - val_loss: 0.3070 - val_accuracy: 0.8994
Learning rate: 1.0000000000000002e-06
Epoch 61/80
467/468 [=====>.] - ETA: 0s - loss: 0.4282 - accuracy: 0.8551
Epoch 00061: ReduceLROnPlateau reducing learning rate to 9.99999974752428e-08.
468/468 [=====] - 219s 468ms/step - loss: 0.4279 - accuracy: 0.8552 - val_loss: 0.3065 - val_accuracy: 0.9006

Performance Analysis

```
In [11]: %load_ext tensorboard
```

```
In [13]: %tensorboard --logdir Logs/
```

Reusing TensorBoard on port 6006 (pid 15592), started 0:22:37 ago. (Use '!kill 15592' to kill it.)



